

SALES_DATA_1

StoreLocation	Product	Date	Revenue
Bellandur	Bananas	January 18,2016	8,236.33
Bellandur	Nutella	January 18,2016	7,455.67
Bellandur	Peanut Butter	January 18,2016	5,316.89
Bellandur	Milk	January 18,2016	2,433.76
Koramangala	Bananas	January 18,2016	9,456.01
Koramangala	Nutella	January 18,2016	3,644.33
Koramangala	Peanut Butter	January 18,2016	8,988.64
Koramangala	Milk	January 18,2016	1,621.58
Bellandur	Bananas	January 17,2016	2342.33
Bellandur	Nutella	January 17,2016	6345.10
Bellandur	Peanut Butter	January 17,2016	5673.01
Bellandur	Milk	January 17,2016	4543.98
Koramangala	Bananas	January 17,2016	8902.65
Koramangala	Nutella	January 17,2016	9114.67
Koramangala	Peanut Butter	January 17,2016	5102.05
Koramangala	Milk	January 17,2016	1299.45

SALES_DATA_2

StoreLocation	Product	Date	Revenue
Bellandur	Bananas	January 18,2016	8,236.33
Bellandur	Nutella	January 18,2016	7,455.67
Bellandur	Peanut Butter	January 18,2016	5,316.89
Bellandur	Milk	January 18,2016	2,433.76
Koramangala	Bananas	January 18,2016	9,456.01
Koramangala	Nutella	January 18,2016	3,644.33
Koramangala	Peanut Butter	January 18,2016	8,988.64
Koramangala	Milk	January 18,2016	1,621.58
Bellandur	Bananas	January 17,2016	2342.33
Bellandur	Nutella	January 17,2016	6345.10
Bellandur	Peanut Butter	January 17,2016	5673.01
Bellandur	Milk	January 17,2016	4543.98
Koramangala	Bananas	January 17,2016	8902.65
Koramangala	Nutella	January 17,2016	9114.67
Koramangala	Peanut Butter	January 17,2016	5102.05
Koramangala	Milk	January 17,2016	1299.45

QUERY_1

```
SELECT Date,Sum(Revenue)
FROM Sales_Data_1
GROUP BY Date;
```

QUERY_2

```
SELECT Date,Sum(Revenue)
FROM Sales_Data_2
GROUP BY Date;
```

SALES_DATA_1

StoreLocation	Product	Date	Revenue
Bellandur	Bananas	January 18,2016	8,236.33
Bellandur	Nutella	January 18,2016	7,455.67
Bellandur	Peanut Butter	January 18,2016	5,316.89
Bellandur	Milk	January 18,2016	2,433.76

SALES_DATA_2

StoreLocation	Product	Date	Revenue
Bellandur	Bananas	January 18,2016	8,236.33
Bellandur	Nutella	January 18,2016	7,455.67
Bellandur	Peanut Butter	January 18,2016	5,316.89
Bellandur	Milk	January 18,2016	2,433.76

IDENTICAL QUERIES ON IDENTICAL DATA

QUERY_1

```
SELECT Date,Sum(Revenue)
FROM Sales_Data_1
GROUP BY Date;
```

SALES_DATA_1

StoreLocation	Product	Date	Revenue
Bellandur	Nutella	January 18,2016	7,455.67
Bellandur	Peanut Butter	January 18,2016	5,316.89
Bellandur	Milk	January 18,2016	2,433.76

QUERY_2

```
SELECT Date,Sum(Revenue)
FROM Sales_Data_2
GROUP BY Date;
```

SALES_DATA_2

StoreLocation	Product	Date	Revenue
Bellandur	Nutella	January 18,2016	7,455.67
Bellandur	Peanut Butter	January 18,2016	5,316.89
Bellandur	Milk	January 18,2016	2,433.76

IDENTICAL QUERIES ON IDENTICAL DATA

QUERY_1

```
SELECT Date, Sum (Revenue)  
FROM Sales_Data_1  
GROUP BY Date;
```

QUERY_2

```
SELECT Date, Sum (Revenue)  
FROM Sales_Data_2  
GROUP BY Date;
```

**BUT QUERY 1 RUNS 10X TO
100X FASTER THAN QUERY 2**

IDENTICAL QUERIES ON
IDENTICAL DATA

BUT QUERY 1 RUNS 10X TO
100X FASTER THAN QUERY 2

THE DIFFERENCE IS
BECAUSE OF AN INDEX

THE DIFFERENCE IS BECAUSE OF AN INDEX

SALES_DATA_1

StoreLocation	Product	Date	Revenue
Bellandur	Bananas	January 18,2016	8,236.33
Bellandur	Nutella	January 18,2016	7,455.67
Bellandur	Peanut Butter	January 18,2016	5,316.89
Bellandur	Milk	January 18,2016	2,433.76
Koramangala	Bananas	January 18,2016	9,456.01
Koramangala	Nutella	January 18,2016	3,644.33
Koramangala	Peanut Butter	January 18,2016	8,988.64
Koramangala	Milk	January 18,2016	1,621.58
Bellandur	Bananas	January 17,2016	2342.33
Bellandur	Nutella	January 17,2016	6345.10
Bellandur	Peanut Butter	January 17,2016	5673.01
Bellandur	Milk	January 17,2016	4543.98
Koramangala	Bananas	January 17,2016	8902.65
Koramangala	Nutella	January 17,2016	9114.67
Koramangala	Peanut Butter	January 17,2016	5102.05
Koramangala	Milk	January 17,2016	1299.45

SALES_DATA_2

StoreLocation	Product	Date	Revenue
Bellandur	Bananas	January 18,2016	8,236.33
Bellandur	Nutella	January 18,2016	7,455.67
Bellandur	Peanut Butter	January 18,2016	5,316.89
Bellandur	Milk	January 18,2016	2,433.76
Koramangala	Bananas	January 18,2016	9,456.01
Koramangala	Nutella	January 18,2016	3,644.33
Koramangala	Peanut Butter	January 18,2016	8,988.64
Koramangala	Milk	January 18,2016	1,621.58
Bellandur	Bananas	January 17,2016	2342.33
Bellandur	Nutella	January 17,2016	6345.10
Bellandur	Peanut Butter	January 17,2016	5673.01
Bellandur	Milk	January 17,2016	4543.98
Koramangala	Bananas	January 17,2016	8902.65
Koramangala	Nutella	January 17,2016	9114.67
Koramangala	Peanut Butter	January 17,2016	5102.05
Koramangala	Milk	January 17,2016	1299.45

THE DIFFERENCE IS BECAUSE OF AN INDEX

THESE TABLES HOLD THE SAME DATA,
BUT BEHIND THE SCENES (ON DISK)
THEY ARE LAID OUT VERY DIFFERENTLY

SALES_DATA_1

StoreLocation	Product	Date	Revenue
Bellandur	Bananas	January 18,2016	8,236.33
Bellandur	Nutella	January 18,2016	7,455.67

SALES_DATA_2

StoreLocation	Product	Date	Revenue
Bellandur	Bananas	January 18,2016	8,236.33
Bellandur	Nutella	January 18,2016	7,455.67

THESE TABLES HOLD THE SAME DATA, BUT BEHIND THE SCENES (ON DISK) THEY ARE LAID OUT VERY DIFFERENTLY

THE DBA INDEXED
SALES_DATA_1
BY DATE

THE DBA FORGOT TO INDEX
SALES_DATA_2 OR
CREATE A PRIMARY KEY

SALES_DATA_1

StoreLocation	Product	Date	Revenue
Bellandur	Bananas	January 18,2016	8,236.33
Bellandur	Nutella	January 18,2016	7,455.67

SALES_DATA_2

StoreLocation	Product	Date	Revenue
Bellandur	Bananas	January 18,2016	8,236.33
Bellandur	Nutella	January 18,2016	7,455.67


THE DBA INDEXED SALES_DATA_1 BY DATE

CONTIGUOUS
LOCATIONS ON DISK



StoreLocation	Product	Date	Revenue
Bellandur	Bananas	January 18,2016	8,236.33
Bellandur	Nutella	January 18,2016	7,455.67
Bellandur	Peanut Butter	January 18,2016	5,316.89
Bellandur	Milk	January 18,2016	2,433.76
Koramangala	Bananas	January 18,2016	9,456.01
Koramangala	Nutella	January 18,2016	3,644.33
Koramangala	Peanut Butter	January 18,2016	8,988.64
Koramangala	Milk	January 18,2016	1,621.58
Bellandur	Bananas	January 17,2016	2342.33

CONTIGUOUS
LOCATIONS ON DISK



Date
January 17,2016
January 18,2016

ACTUAL DISK LAYOUT OF INDEX ON
COLUMN DATE

StoreLocation	Product	Date	Revenue
Bellandur	Bananas	January 17,2016	2342.33
Bellandur	Nutella	January 17,2016	6345.10
Bellandur	Peanut Butter	January 17,2016	5673.01
Bellandur	Milk	January 17,2016	4543.98
Koramangala	Bananas	January 17,2016	8902.65
Koramangala	Nutella	January 17,2016	9114.67
Koramangala	Peanut Butter	January 17,2016	5102.05
Koramangala	Milk	January 17,2016	1299.45

CONTIGUOUS
LOCATIONS ON DISK




SALES_DATA_1

StoreLocation	Product	Date	Revenue
CONCEPTUAL LAYOUT			


THE DBA INDEXED SALES_DATA_1 BY DATE

CONTIGUOUS
LOCATIONS ON DISK



StoreLocation	Product	Date	Revenue
Bellandur	Bananas	January 18,2016	8,236.33
Bellandur	Nutella	January 18,2016	7,455.67
Bellandur	Peanut Butter	January 18,2016	5,316.89
Bellandur	Milk	January 18,2016	2,433.76
Koramangala	Bananas	January 18,2016	9,456.01
Koramangala	Nutella	January 18,2016	3,644.33
Koramangala	Peanut Butter	January 18,2016	8,988.64
			1,621.58
			2342.33

CONTIGUOUS
LOCATIONS ON DISK




```
SELECT Date , Sum (Revenue)
FROM Sales_Data_1
GROUP BY Date;
```

SALES_DATA_1

StoreLocation	Product	Date	Revenue
CONCEPTUAL LAYOUT			

CONTIGUOUS
LOCATIONS ON DISK



			Revenue
			2342.33
			6345.10
	Butter	January 17,2016	5673.01
		January 17,2016	4543.98
	as	January 17,2016	8902.65
Koramangala	Nutella	January 17,2016	9114.67
Koramangala	Peanut Butter	January 17,2016	5102.05
Koramangala	Milk	January 17,2016	1299.45

NOW A QUERY
LIKE THIS IS
EASY-PEASY!

```
SELECT Date, Sum (Revenue)  
FROM Sales_Data_1  
GROUP BY Date;
```

THESE TABLES HOLD THE SAME DATA, BUT BEHIND THE SCENES (ON DISK) THEY ARE LAID OUT VERY DIFFERENTLY

THE DBA INDEXED
SALES_DATA_1
BY DATE

THE DBA FORGOT TO INDEX
SALES_DATA_2 OR
CREATE A PRIMARY KEY

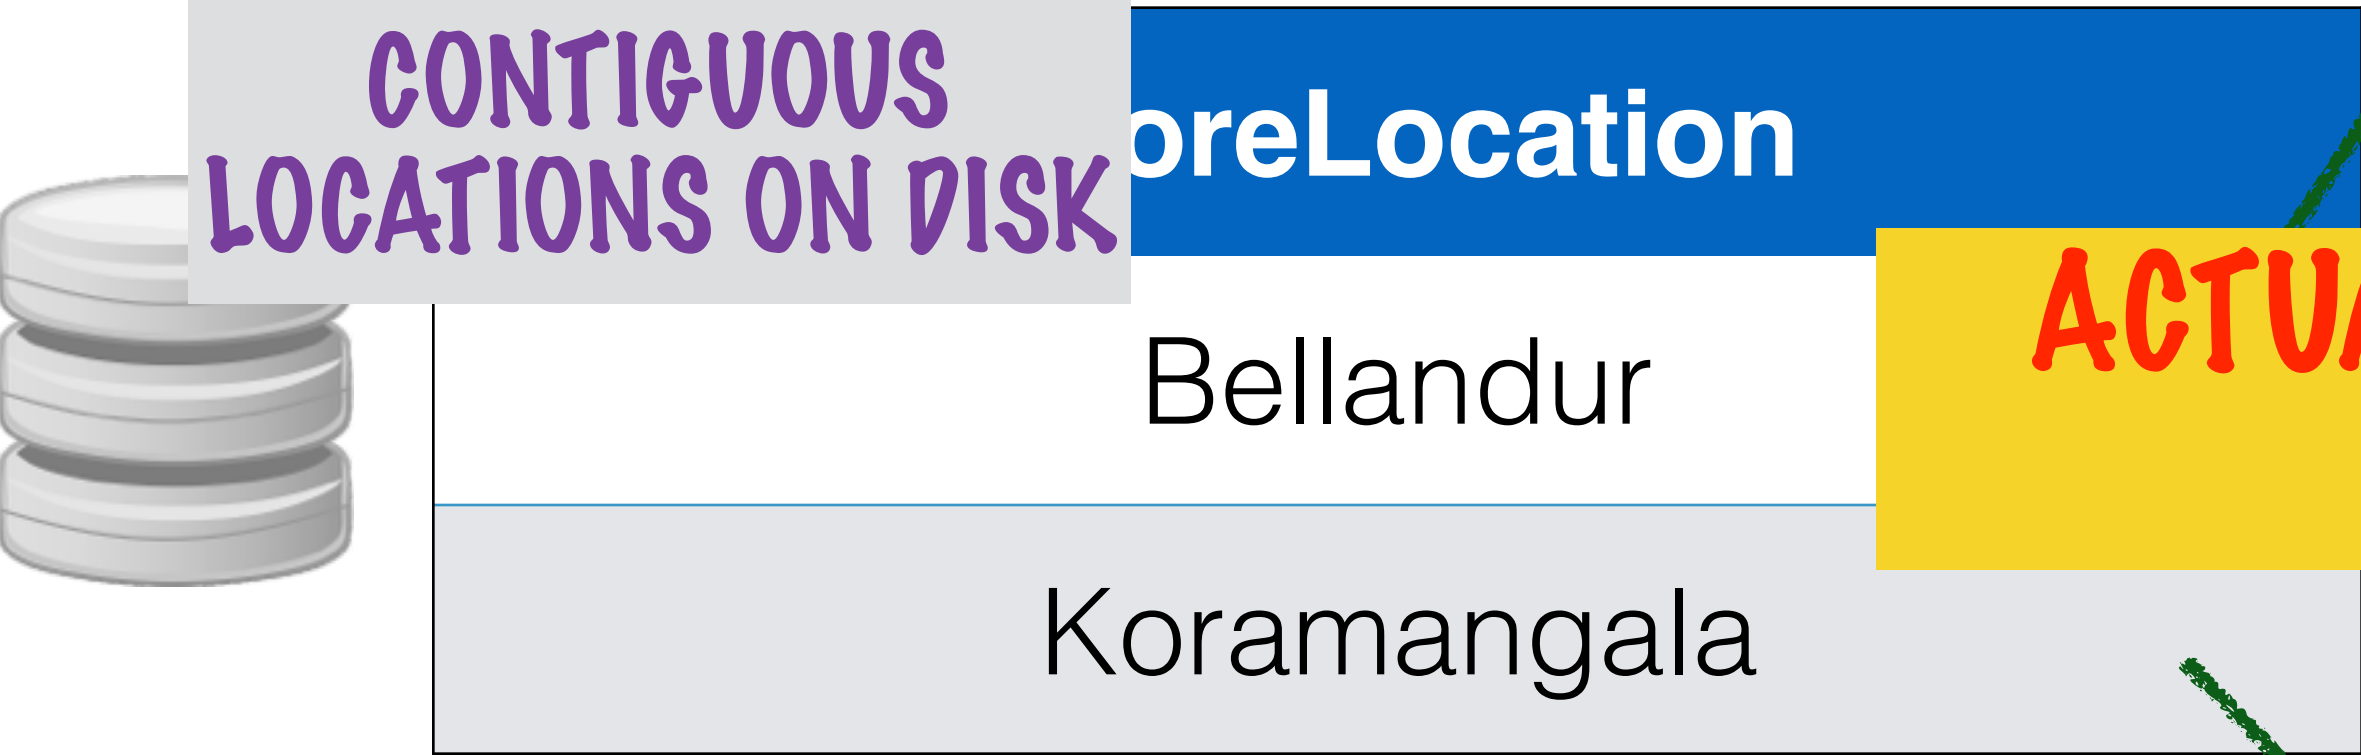
SALES_DATA_1


StoreLocation	Product	Date	Revenue
Bellandur	Bananas	January 18,2016	8,236.33
Bellandur	Nutella	January 18,2016	7,455.67

SALES_DATA_2

StoreLocation	Product	Date	Revenue
Bellandur	Bananas	January 18,2016	8,236.33
Bellandur	Nutella	January 18,2016	7,455.67

THE DBA FORGOT TO INDEX SALES DATA 2 OR CREATE A PRIMARY KEY



CONTIGUOUS LOCATIONS ON DISK		Product	Date	Revenue
	Bellandur	Bananas	January 18,2016	8,236.33
		Nutella	January 18,2016	7,455.67
	Bellandur	Peanut Butter	January 18,2016	5,316.89
	Bellandur	Milk	January 18,2016	2,433.76
	Bellandur	Bananas	January 17,2016	2342.33
	Bellandur	Bananas	January 17,2016	2342.33
	Bellandur	Nutella	January 17,2016	6345.10
	Bellandur	Peanut Butter	January 17,2016	5673.01
	Bellandur	Milk	January 17,2016	4543.98

ACTUAL DISK LAYOUT, PICKED BY THE DBMS "AT RANDOM"


SALES_DATA_2

StoreLocation	Product	Date	Revenue
CONCEPTUAL LAYOUT			

StoreLocation	Product	Date	Revenue
Koramangala	Bananas	January 18,2016	9,456.01
Koramangala	Nutella	January 18,2016	3,644.33
Koramangala	Peanut Butter	January 18,2016	8,988.64
	Milk	January 18,2016	1,621.58
Koramangala	Bananas	January 17,2016	8902.65
Koramangala	Nutella	January 17,2016	9114.67
Koramangala	Peanut Butter	January 17,2016	5102.05
Koramangala	Milk	January 17,2016	1299.45


THE DBA FORGOT TO INDEX
SALES DATA 2 OR
CREATE A PRIMARY KEY

CONTIGUOUS
LOCATIONS ON DISK



Product		Date	Revenue
Bananas		January 18,2016	8,236.33
Nutella		January 18,2016	7,455.67
Bellandur	Peanut Butter	January 18,2016	5,316.89
Bellandur	Milk	January 18,2016	2,433.76
Bellandur	Bananas	January 17,2016	2342.33
Bellandur	Bananas	January 17,2016	2342.33
Bellandur	Nutella	January 17,2016	6345.10
Bellandur	Peanut Butter	January 17,2016	5673.01
Bellandur	Milk	January 17,2016	4543.98

CONTIGUOUS
LOCATIONS ON DISK




```
SELECT Date , Sum (Revenue)
FROM Sales_Data_2
GROUP BY Date;
```

ACROSS
VDBMS AT RANDOM

SALES_DATA_2

StoreLocation	Product	Date	Revenue
CONCEPTUAL LAYOUT			

CONTIGUOUS
LOCATIONS ON DISK



			Revenue
2016			9,456.01
2016			3,644.33
Peanut Butter		January 18,2016	8,988.64
Milk		January 18,2016	1,621.58
Bananas		January 17,2016	8902.65
Koramangala	Nutella	January 17,2016	9114.67
Koramangala	Peanut Butter	January 17,2016	5102.05
Koramangala	Milk	January 17,2016	1299.45

01 01 01..THIS
QUERY NEEDS A
TON OF DISK IO!

```
SELECT Date, Sum (Revenue)  
FROM Sales_Data_2  
GROUP BY Date;
```

MORAL OF THE STORY:

**SETTING UP INDICES CORRECTLY
CAN SAVE HOURS AND HOURS
OF TIME**

OKEY-DOKEY, SO HOW
DO I SET UP INDICES
CORRECTLY?



OKEY-DOKEY, SO HOW DO I
SET UP INDICES CORRECTLY?

```
CREATE INDEX IX_Date ON Sales_Data_1(Date) ;
```

THE SQL IS
SIMPLE ENOUGH

OKEY-DOKEY, SO HOW DO I
SET UP INDICES CORRECTLY?

```
CREATE INDEX IX_Date ON Sales_Data_1 (Date) ;
```

THE SQL IS
SIMPLE ENOUGH

OKEY-DOKEY, SO HOW DO I
SET UP INDICES CORRECTLY?

```
CREATE INDEX IX_Date ON Sales_Data_1 (Date) ;  
          (INDEX NAME)
```

THE SQL IS
SIMPLE ENOUGH

OKEY-DOKEY, SO HOW DO I
SET UP INDICES CORRECTLY?

```
CREATE INDEX IX_Date ON Sales_Data_1 (Date) ;
```

(TABLE NAME)

THE SQL IS
SIMPLE ENOUGH

OKEY-DOKEY, SO HOW DO I
SET UP INDICES CORRECTLY?

```
CREATE INDEX IX_Date ON Sales_Data_1(Date) ;  
                                     (COLUMNS)
```

THE SQL IS
SIMPLE ENOUGH

OKEY-DOKEY, SO HOW DO I SET UP INDICES CORRECTLY?

```
CREATE INDEX IX_Store_Product ON  
Sales_Data_1 (StoreLocation, ProductName) ;
```

(MULTIPLE COLUMNS
ARE FINE TOO)

THE SQL IS
SIMPLE ENOUGH

OKEY-DOKEY, SO HOW DO I
SET UP INDICES CORRECTLY?

YOU CAN USE **ALTER** TO CHANGE
AN EXISTING INDEX, AND **DROP**
TO GET RID OF AN INDEX

THE SQL IS
SIMPLE ENOUGH

OKEY-DOKEY, SO HOW DO I
SET UP INDICES CORRECTLY?

INDICES MAKE QUERIES MUCH
FASTER, BUT UPDATES MUCH SLOWER

A CLASSIC DATABASES
TRADE-OFF

**INDICES MAKE QUERIES MUCH
FASTER, BUT UPDATES MUCH SLOWER**

**MOST DBMS WILL - BY DEFAULT-
CREATE INDICES ON PRIMARY KEYS**

**INDICES MAKE QUERIES MUCH
FASTER, BUT UPDATES MUCH SLOWER**

**MOST DBMS WILL - BY DEFAULT- CREATE INDICES ON PRIMARY KEYS
INDICES ON COMMONLY QUERIED
COLUMNS ARE A LIFE-SAVER**

INDICES MAKE LESS SENSE ON COLUMNS THAT

**ARE IN VERY
SMALL TABLES**

**HOLD A LOT OF
NULLS**

**ARE WRITTEN
MORE THAN READ**